

Branch: master ▼

Find file

Copy path

Multi-class-with-imbalanced-dataset-classification / 20-news-group-classification.ipynb

 javaidnabi31 Update 20-news-group-classification.ipynb

d0021b4 on 2 Jan

1 contributor



Raw

Blame

History



1128 lines (1127 sloc) 421 KB

Machine Learning- Multi-class Classification with Imbalanced Data-set

In [33]:

```
import pandas as pd
import numpy as np
import pickle
from keras.preprocessing.text import Tokenizer
from keras.models import Sequential, Model
from keras.layers import Activation, Dense, Dropout
from sklearn.preprocessing import LabelBinarizer
import sklearn.datasets as skds
from pathlib import Path
import matplotlib.pyplot as plt
import itertools
from sklearn.metrics import confusion_matrix
# For reproducibility
np.random.seed(1237)
```

In [34]:

```
# Source file directory
# downlaod from http://qwone.com/~jason/20Newsgroups/

path_train = "20news-bydate/20news-bydate-train"

files_train = skds.load_files(path_train, load_content=False)

label_index = files_train.target
label_names = files_train.target_names
labelled_files = files_train filenames

data_tags = ["filename", "category", "news"]
data_list = []

# Read and add data from file to a list
i=0
for f in labelled_files:
    data_list.append((f, label_names[label_index[i]], Path(f).read_text()))
    i += 1

# We have training data available as dictionary filename, category, data
data = pd.DataFrame.from_records(data_list, columns=data_tags)

# 20 news groups
num_labels = 20
vocab_size = 15000
batch_size = 100
```

```

num_epochs = 30

# lets take 80% data as training and remaining 20% for tes
t.
train_size = int(len(data) * .8)

train_posts = data['news'][:train_size]
train_tags = data['category'][:train_size]
train_files_names = data['filename'][:train_size]

test_posts = data['news'][train_size:]
test_tags = data['category'][train_size:]
test_files_names = data['filename'][train_size:]

# define Tokenizer with Vocab Size
tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train_posts)

x_train = tokenizer.texts_to_matrix(train_posts, mode='tfidf')
x_test = tokenizer.texts_to_matrix(test_posts, mode='tfidf')

encoder = LabelBinarizer()
encoder.fit(train_tags)
y_train = encoder.transform(train_tags)
y_test = encoder.transform(test_tags)

```

In [35]:

```

#crate class imbalance; my rough approach
data_imb = data.copy()

for i in range(1,6):
    for index, row in data_imb.iterrows():
        if((row["category"] == "alt.atheism" or row["category"] == "talk.politics.misc"
            or row["category"] == "soc.religion.christian"
or row["category"] == "talk.politics.mideast")and (index % 3) == 0):
            data_imb.drop(index, inplace=True)
        elif((row["category"] == "comp.os.ms-windows.misc" or row["category"] == "comp.sys.ibm.pc.hardware" or row["category"] == "comp.graphics"
            or row["category"] == "comp.windows.x" or row["category"] == "comp.sys.mac.hardware")and (index % 4) == 0):
            data_imb.drop(index, inplace=True)
        elif((row["category"] == "sci.med" or row["category"] == "sci.space" or row["category"] == "sci.electronics"
            or row["category"] == "sci.crypt" or row["category"] == "misc.forsale") and (index % 5) == 0):
            data_imb.drop(index, inplace=True)

data_imb.reset_index(drop = True, inplace=True)

```

In [36]:

```
data_imb.category.value_counts()
```

Out[36]:

```
rec.sport.hockey          600
rec.motorcycles           598
rec.sport.baseball        597
rec.autos                 594
talk.politics.guns        546
talk.religion.misc        377
sci.med                   207
sci.electronics           205
sci.space                 197
sci.crypt                 183
misc.forsale              171
comp.os.ms-windows.misc  151
comp.graphics             146
comp.sys.ibm.pc.hardware  137
comp.windows.x            136
comp.sys.mac.hardware     131
soc.religion.christian     86
talk.politics.mideast     67
alt.atheism               63
talk.politics.misc        55
Name: category, dtype: int64
```

In [37]:

```
class_labels = data_imb.category.tolist()
```

In [38]:

```
data_imb.head()
```

Out[38]:

	filename	category	news
0	20news-bydate/20news-bydate-train\rec.sport.ba...	rec.sport.baseball	From: cubbie@...
1	20news-bydate/20news-bydate-train\comp.sys.mac...	comp.sys.mac.hardware	From: gnelson(Gregory Nelsc
2	20news-bydate/20news-bydate-train\sci.crypt\15246	sci.crypt	From: crypt-comments@m... C...
3	20news-bydate/20news-bydate-	comp.sys.mac.hardware	From: ()\nSubje Problems??...

	train\comp.sys.mac...		
4	20news- bydate/20news- bydate- train\alt.atheism\...	alt.atheism	From: keith@c Allan Schne...

In [111]:

```
# 20 news groups
num_labels = 20
vocab_size = 15000
batch_size = 100
num_epochs = 30

# lets take 80% data as training and remaining 20% for test.
train_size = int(len(data_imb) * .8)

train_posts = data_imb['news'][:train_size]
train_tags = data_imb['category'][:train_size]
train_files_names = data_imb['filename'][:train_size]

test_posts = data_imb['news'][train_size:]
test_tags = data_imb['category'][train_size:]
test_files_names = data_imb['filename'][train_size:]

# define Tokenizer with Vocab Size
tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train_posts)

x_train = tokenizer.texts_to_matrix(train_posts)
x_test = tokenizer.texts_to_matrix(test_posts)

encoder = LabelBinarizer()
encoder.fit(train_tags)
y_train = encoder.transform(train_tags)
y_test = encoder.transform(test_tags)
```

In [112]:

```
x_train.shape, y_train.shape
```

Out[112]:

```
((4197, 15000), (4197, 20))
```

In [113]:

```
#let us build a basic model
model = Sequential()
model.add(Dense(512, input_shape=(vocab_size,)))
model.add(Activation('relu'))
model.add(Dropout(0.3))
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.3))
```

```

model.add(Dropout(0.5))
model.add(Dense(num_labels))
model.add(Activation('softmax'))
model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

```

Layer (type) am #	Output Shape	Par
dense_22 (Dense) 0512	(None, 512)	768
activation_16 (Activation)	(None, 512)	0
dropout_11 (Dropout)	(None, 512)	0
dense_23 (Dense) 656	(None, 512)	262
activation_17 (Activation)	(None, 512)	0
dropout_12 (Dropout)	(None, 512)	0
dense_24 (Dense) 60	(None, 20)	102
activation_18 (Activation)	(None, 20)	0
Total params: 7,953,428		
Trainable params: 7,953,428		
Non-trainable params: 0		

In [114]:

```

num_epochs = 10
batch_size = 128
history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=num_epochs,
                    verbose=2,
                    validation_split=0.2)

```

Train on 3357 samples, validate on 840 samples

```

Epoch 1/10
  - 6s - loss: 2.1494 - acc: 0.4242 - val_loss: 1.1093 - va
l_acc: 0.7298
Epoch 2/10
  - 2s - loss: 0.5979 - acc: 0.8642 - val_loss: 0.5704 - va
l_acc: 0.8429
Epoch 3/10
  - 2s - loss: 0.1353 - acc: 0.9762 - val_loss: 0.3956 - va
l_acc: 0.8798
Epoch 4/10
  - 2s - loss: 0.0333 - acc: 0.9973 - val_loss: 0.3606 - va
l_acc: 0.8976
Epoch 5/10
  - 2s - loss: 0.0112 - acc: 1.0000 - val_loss: 0.3540 - va
l_acc: 0.8952
Epoch 6/10
  - 2s - loss: 0.0064 - acc: 0.9997 - val_loss: 0.3522 - va
l_acc: 0.8893
Epoch 7/10
  - 2s - loss: 0.0035 - acc: 1.0000 - val_loss: 0.3509 - va
l_acc: 0.8905
Epoch 8/10
  - 2s - loss: 0.0026 - acc: 1.0000 - val_loss: 0.3551 - va
l_acc: 0.8893
Epoch 9/10
  - 2s - loss: 0.0019 - acc: 1.0000 - val_loss: 0.3525 - va
l_acc: 0.8893
Epoch 10/10
  - 2s - loss: 0.0017 - acc: 1.0000 - val_loss: 0.3600 - va
l_acc: 0.8857

```

In []:

```

score, acc = model.evaluate(x_test, y_test,
                             batch_size=batch_size, verbose=2)

print('Test accuracy:', acc)

```

In [27]:

```

#another approach using GRU model, takes longer time
from tensorflow.python.keras.preprocessing.text imp
ort Tokenizer
from tensorflow.python.keras.preprocessing.sequence
import pad_sequences

tokenizer_obj = Tokenizer()

tokenizer_obj.fit_on_texts(train_posts)

# pad sequences
max_length = max([len(s.split()) for s in train_posts])

# define vocabulary size
vocab_size = len(tokenizer_obj.word_index) + 1

X_train_tokens = tokenizer_obj.texts_to_sequences(train_p
osts)

```

```

X_test_tokens = tokenizer_obj.texts_to_sequences(test_posts)

X_train_pad = pad_sequences(X_train_tokens, maxlen=max_length, padding='post')
X_test_pad = pad_sequences(X_test_tokens, maxlen=max_length, padding='post')

encoder = LabelBinarizer()
encoder.fit(train_tags)
y_train = encoder.transform(train_tags)
y_test = encoder.transform(test_tags)

```

In [31]:

```

#another approach using GRU model, takes longer time
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, GRU
from keras.layers.embeddings import Embedding

EMBEDDING_DIM = 100

print('Build model...')

model = Sequential()
model.add(Embedding(vocab_size, EMBEDDING_DIM, input_length=max_length))
model.add(GRU(units=32, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(num_labels, activation='softmax'))

# try using different optimizers and different optimizer configs
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

print('Summary of the built model...')
print(model.summary())

```

Build model...

Summary of the built model...

Layer (type) am #	Output Shape	Param #
=====		
embedding_6 (Embedding)	(None, 50, 50)	3599550
=====		
gru_6 (GRU)	(None, 32)	7968
=====		
dense_6 (Dense)	(None, 20)	660
=====		


```
=====
Total params: 3,608,178
Trainable params: 3,608,178
Non-trainable params: 0
```

None

In [116]:

```
text_labels = encoder.classes_

for i in range(10):
    prediction = model.predict(np.array([x_test[i]]))
    predicted_label = text_labels[np.argmax(prediction[0
])]
    #print(test_files_names.iloc[i])
    print('Actual label:' + test_tags.iloc[i])
    print("Predicted label: " + predicted_label)

def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`
    e`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.new
axis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    # print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), ran
ge(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else
"black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```


```
y_pred = model.predict(x_test);
cnf_matrix = confusion_matrix(np.argmax(y_test, axis=1), n
p.argmax(y_pred, axis=1))
```


```
# Plot normalized confusion matrix
fig = plt.figure()
fig.set_size_inches(14, 12, forward=True)
#fig.align_labels()

# fig.subplots_adjust(left=0.0, right=1.0, bottom=0.0, top
=1.0)
plot_confusion_matrix(cnf_matrix, classes=np.asarray(label
_names), normalize=True,
                      title='Normalized confusion matrix')

fig.savefig("txt_classification-smote" + str(num_epochs) +
".png", pad_inches=5.0)
```

Actual label:sci.med
Predicted label: sci.med
Actual label:sci.crypt
Predicted label: sci.crypt
Actual label:rec.motorcycles
Predicted label: rec.motorcycles
Actual label:comp.sys.mac.hardware
Predicted label: sci.electronics
Actual label:talk.politics.guns
Predicted label: talk.politics.guns
Actual label:soc.religion.christian
Predicted label: soc.religion.christian
Actual label:comp.graphics
Predicted label: comp.graphics
Actual label:rec.sport.baseball
Predicted label: rec.sport.baseball
Actual label:talk.politics.guns
Predicted label: talk.politics.guns
Actual label:rec.motorcycles
Predicted label: rec.motorcycles
Normalized confusion matrix





In [117]:

```
prediction = model.predict(x_test)
```

In [118]:

```
predictions = np.argmax(prediction, axis = 1)
```

In [119]:

```
predictions
```

Out[119]:

```
array([13, 11,  8, ...,  8,  2, 16], dtype=int64)
```

In [120]:

```
y_test_labels = np.argmax(y_test, axis =1)
```

In [121]:

```
y_test_labels
```

Out[121]:

```
array([13, 11,  8, ...,  8,  5, 16], dtype=int64)
```

In [93]:

```
y_train_labels = np.argmax(y_train, axis =1)
y_train_labels
```

Out[93]:

```
array([ 9,  4, 11, ...,  9, 14, 16], dtype=int64)
```

In [94]:

```
#The Kappa score tell you how much better, or worse, your
  classifier is than what would be expected by random chanc
  e.
  #If you were to randomly assign cases to classes (i.e. a k
  ind of terribly uninformed classifier), you'd get some cor
  rect simply by chance.
  #Therefore, you will always find that the Kappa value is l
  ower than the overall accuracy.
  #The Kappa index is however considered to be a more conser
  vative measure than the overall classification accuracy.
  #Your KIA value is telling you essentially that your class
  ifier is about 66% better than a random assignment of case
  s to the various classes. That's not bad!
  #A kappa value of 1 represents perfect agreement, while a
  value of 0 represents no agreement
```

value of 0 represents no agreement.

```
from sklearn.metrics import cohen_kappa_score
cohen_score = cohen_kappa_score(y_test_labels, predictions
)
```

In [175]:

```
cohen_score
```

Out[175]:

0.8570501222325042

In [55]:

```
from sklearn.metrics import precision_recall_fscore_support as score
```

```
precision, recall, fscore, support = score(y_test_labels,
predictions)
```

```
print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))
```

```
precision: [0.75      0.66666667 0.75      0.68965517 0.
7037037  0.82142857
 0.9      0.89719626 0.890625   0.94308943 0.97637795 0.
96774194
 0.88235294 0.92105263 0.875      1.      0.87826087 1.
 0.76923077 0.91463415]
recall: [0.5      0.75862069 0.82758621 0.64516129 0.76
0.88461538
 0.77142857 0.88888889 0.97435897 0.95867769 1.      0.
96774194
 0.66666667 0.81395349 0.92105263 0.65      0.98058252 0.
85714286
 0.55555556 0.92592593]
fscore: [0.6      0.70967742 0.78688525 0.66666667 0.730
76923 0.85185185
 0.83076923 0.89302326 0.93061224 0.95081967 0.98804781 0.
96774194
 0.75949367 0.86419753 0.8974359  0.78787879 0.9266055  0.
92307692
 0.64516129 0.9202454 ]
support: [ 12  29  29  31  25  26  35 108 117 121 124  31
 45  43  38  20 103  14
 18  81]
```

In [58]:

```
from sklearn.metrics import confusion_matrix
import numpy as np
```

```
cm = confusion_matrix(y_test_labels, predictions)
recall = np.diag(cm) / np.sum(cm, axis = 1)
```

```
precision = np.diag(cm) / np.sum(cm, axis = 0)

print(recall)

print(precision)
```

```
[0.5          0.75862069 0.82758621 0.64516129 0.76          0.
88461538
 0.77142857 0.88888889 0.97435897 0.95867769 1.          0.
96774194
 0.66666667 0.81395349 0.92105263 0.65          0.98058252 0.
85714286
 0.55555556 0.92592593]
[0.75          0.66666667 0.75          0.68965517 0.7037037 0.
82142857
 0.9          0.89719626 0.890625     0.94308943 0.97637795 0.
96774194
 0.88235294 0.92105263 0.875          1.          0.87826087 1.
 0.76923077 0.91463415]
```

In []:

```
#Let us try some sampling technique to remove class imbalance
from imblearn.over_sampling import SMOTE
#Over-sampling: SMOTE
#SMOTE (Synthetic Minority Oversampling TEchnique) consist
s of synthesizing elements for the minority class,
#based on those that already exist. It works randomly pick
ing a point from the minority class and computing
#the k-nearest neighbors for this point. The synthetic poin
ts are added between the chosen point and its neighbors.
#We'll use ratio='minority' to resample the minority clas
s.
smote = SMOTE('minority')

X_sm, y_sm = smote.fit_sample(x_train, y_train)
print(X_sm.shape, y_sm.shape)
```

In [126]:

```
from sklearn.utils import class_weight
class_weight = class_weight.compute_class_weight('balance
d', np.unique(y_train_labels), y_train_labels)
num_epochs = 10
batch_size = 128
history = model.fit(X_sm, y_sm,
                    batch_size=batch_size,
                    epochs=num_epochs,
                    verbose=2,
                    class_weight=class_weight,
                    validation_split=0.2)
```

Train on 3716 samples, validate on 930 samples

Epoch 1/10

- 10s - loss: 0.0593 - acc: 0.9839 - val_loss: 0.2841 - v
al_acc: 0.9075

Epoch 2/10

0.011 0.998 0.012 0.997 1.1 0.1010

```

- 2s - loss: 0.0138 - acc: 0.9995 - val_loss: 0.1916 - va
l_acc: 0.9441
Epoch 3/10
- 3s - loss: 0.0068 - acc: 0.9997 - val_loss: 0.1903 - va
l_acc: 0.9387
Epoch 4/10
- 3s - loss: 0.0057 - acc: 0.9997 - val_loss: 0.1924 - va
l_acc: 0.9376
Epoch 5/10
- 2s - loss: 0.0054 - acc: 0.9997 - val_loss: 0.1889 - va
l_acc: 0.9452
Epoch 6/10
- 2s - loss: 0.0051 - acc: 0.9997 - val_loss: 0.1899 - va
l_acc: 0.9430
Epoch 7/10
- 3s - loss: 0.0050 - acc: 0.9997 - val_loss: 0.1897 - va
l_acc: 0.9419
Epoch 8/10
- 2s - loss: 0.0048 - acc: 0.9997 - val_loss: 0.1889 - va
l_acc: 0.9409
Epoch 9/10
- 2s - loss: 0.0047 - acc: 0.9997 - val_loss: 0.1900 - va
l_acc: 0.9398
Epoch 10/10
- 2s - loss: 0.0047 - acc: 0.9997 - val_loss: 0.1889 - va
l_acc: 0.9409

```

In [127]:

```

score, acc = model.evaluate(x_test, y_test,
                             batch_size=batch_size, verbose=2)

print('Test accuracy:', acc)

```

Test accuracy: 0.8904761907032558

In [128]:

```

from sklearn.metrics import roc_curve, auc
from scipy import interp
from itertools import cycle

fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(num_labels):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], prediction
[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Compute micro-average ROC curve and ROC area
fpr["micro"], tpr["micro"], _ = roc_curve(y_test.ravel(),
prediction.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

# Compute macro-average ROC curve and ROC area

# First aggregate all false positive rates
all_fpr = np.unique(np.concatenate([fpr[i] for i in range

```

```

all_tpr = np.unique(np.concatenate([tpr[i] for i in range
(num_labels)]))

# Then interpolate all ROC curves at this points
mean_tpr = np.zeros_like(all_fpr)
for i in range(num_labels):
    mean_tpr += interp(all_fpr, fpr[i], tpr[i])

# Finally average it and compute AUC
mean_tpr /= num_labels

```